

Vertica7.2_Vmart_T-SQL_Queries_Yuan_Wang

Tuesday, August 30, 2016 2:39 PM

Introducing the VMart Example Database

Vertica ships with a sample multi-schema database called the VMart Example Database, which represents a database that might be used by a large supermarket (VMart) to access information about its products, customers, employees, and online and physical stores. Using this example, you can create, run, optimize, and test a multi-schema database.

The VMart database contains the following schema:

- public (automatically created in any newly created Vertica database)
- store
- online_Sales

The complete graph of these three schema is available on [Appendix: VMart Example Database Schema](#).

```
localhost login: yuan
Password:
Last login: Tue Aug 30 14:41:12 on tty1
[yuan@localhost ~]$ su - dbadmin
Password:
Last login: Mon Aug 29 16:54:01 EDT 2016 on tty1
[dbadmin@localhost ~]$ vsql
Welcome to vsql, the Vertica Analytic Database interactive terminal.

Type: \h or \? for help with vsql commands
      \g or terminate with semicolon to execute query
      \q to quit

dbadmin=> select * from databases;
 database_id | database_name | owner_id | owner_name | start_time | compliance_message
-----+-----+-----+-----+-----+-----
 45035996273704976 | VMart | 45035996273704962 | dbadmin | 2016-08-29 12:10:34.036495-04 | The database is in compliance with respect to raw data siz
e. | 0 | none
(1 row)
```

```
dbadmin=> select table_name from tables where table_schema='online_sales';
 table_name
-----
online_page_dimension
call_center_dimension
online_sales_fact
(3 rows)
```

```
dbadmin=> select table_name from tables where table_schema='store';
 table_name
-----
store_dimension
store_sales_fact
store_orders_fact
(3 rows)
```

```
dbadmin=> select table_name from tables where table_schema='public';
 table_name
-----
customer_dimension
product_dimension
promotion_dimension
date_dimension
vendor_dimension
employee_dimension
shipping_dimension
warehouse_dimension
inventory_fact
(9 rows)
```

Transcat-SQL Queries

I have used different T-SQL aggregate and analytic functions to explore this database to get the information I am interested in.

Query 1:

The following query lists name and annual income of male and female customer who has the highest annual income.

```
ldbadmin@localhost ~1$ less /mnt/hgfs/UMShare/Max_Annual.sql
SELECT customer_name, annual_income
FROM public.customer_dimension
WHERE (customer_gender, annual_income) IN (
  SELECT customer_gender, MAX(annual_income)
  FROM public.customer_dimension
  GROUP BY customer_gender);
(END)
```

Output:

```
dbadmin=> \i '/mnt/hgfs/UMShare/Max_Annual.sql'
customer_name | annual_income
-----+-----
James M. McNulty | 9999979
Emily G. Vogel | 9999998
(2 rows)
```

Query 2:

The following query calculates the average rent cost of all stores in each state of U.S. and rank it from the highest to lowest.

```
ldbadmin@localhost ~1$ less /mnt/hgfs/UMShare/rank_store.sql
SELECT store_state, round(avg(monthly_rent_cost),2) as state_avg_rent,
RANK() OVER(
  ORDER BY avg(monthly_rent_cost) desc) AS RANK
FROM store.store_dimension
GROUP BY store_state;
(END)
```

Output:

```
dbadmin=> \i /mnt/hgfs/UMShare/rank_store.sql
store_state | state_avg_rent | RANK
-----+-----+-----
NJ | 10042.2 | 1
VA | 9806.5 | 2
MD | 9721 | 3
WI | 8645.67 | 4
NH | 8605.5 | 5
IA | 8406 | 6
SD | 8100.75 | 7
CO | 7244.92 | 8
MS | 7205 | 9
DC | 7077.33 | 10
CT | 6509.82 | 11
MA | 6476.1 | 12
IN | 6462.8 | 13
FL | 6387.91 | 14
IL | 6358.64 | 15
WA | 6341 | 16
KS | 6192.6 | 17
GA | 6129.22 | 18
PA | 5858.11 | 19
TN | 5750.83 | 20
MI | 5548.68 | 21
CA | 5509.66 | 22
TX | 5475.33 | 23
LA | 5004.5 | 24
NU | 4956 | 25
NC | 4875 | 26
NY | 4573 | 27
UT | 4505 | 28
AZ | 3805.25 | 29
OR | 2408 | 30
SC | 2332 | 31
(31 rows)
```

Query 3

The following query calculates the minimum, maximum, average employee annual salary and number of employees of all the stores in the state 'MA'

```
[dbadmin@localhost ~]$ less /mnt/hgfs/UMShare/count.sql
SELECT DISTINCT s.store_name
      , MIN(annual_salary) OVER (PARTITION BY s.store_name) AS MinSalary
      , MAX(annual_salary) OVER (PARTITION BY s.store_name) AS MaxSalary
      , round(AVG(annual_salary) OVER (PARTITION BY s.store_name),1) AS AvgSalary
      ,COUNT(ed.employee_key) OVER (PARTITION BY s.store_name) AS EmployeesPerStore
FROM store.store_dimension AS s
JOIN store.store_orders_fact as sof
  ON s.store_key = sof.store_key
JOIN public.employee_dimension AS ed
  ON sof.employee_key = ed.employee_key
WHERE s.store_state = 'MA'
ORDER BY AvgSalary DESC;
(END)
```

Output:

```
dbadmin=> \i /mnt/hgfs/UMShare/count.sql
store_name | MinSalary | MaxSalary | AvgSalary | EmployeesPerStore
-----+-----+-----+-----+-----
Store138   |      1202 |    199553 |    55765.3 |           1236
Store12    |      1202 |    199992 |    54901.1 |           1248
Store58    |      1202 |    992363 |    54164.3 |           1144
Store210   |      1204 |    199766 |    54033.8 |           1208
Store36    |      1200 |    198786 |    53005.5 |           1198
Store244   |      1202 |    199903 |    52693.1 |           1143
Store75    |      1202 |    199225 |    52379.2 |           1180
Store57    |      1200 |    199955 |    52083.4 |           1206
Store157   |      1202 |    199744 |    51935.8 |           1255
Store237   |      1202 |    199282 |    51119.8 |           1167
(10 rows)
```

Query 4:

The following query finds the profit for that calendar month and returns a running/cumulative average

```
[dbadmin@localhost ~]$ less /mnt/hgfs/UMShare/avg_analytic.sql
SELECT calendar_month_number_in_year, SUM(product_price-product_cost) AS profit,
      ROUND(AVG(SUM(product_price-product_cost))
      OVER (ORDER BY calendar_month_number_in_year),2) AS moving_profit
FROM product_dimension AS p
JOIN inventory_fact AS i ON p.product_key = i.product_key
JOIN date_dimension AS d ON d.date_key = i.date_key
GROUP BY calendar_month_number_in_year;
(END)
```

Output:

```
dbadmin=> \i /mnt/hgfs/UMShare/avg_analytic.sql
calendar_month_number_in_year | profit | moving_profit
-----+-----+-----
1 | 12022340 |      12022340
2 |  9858008 |      10940174
3 | 11496053 |      11125467
4 | 11534402 |     11227700.75
5 | 11933574 |     11368875.4
6 | 11331899 |     11362712.67
7 | 10813047 |     11284189
8 | 12485303 |     11434328.25
9 | 10912531 |     11376350.78
10 | 11883274 |     11427043.1
11 | 10618792 |     11353565.73
12 | 12453314 |     11445211.42
(12 rows)
```